

Mathematical Biostatistics Bootcamp: Lecture 9, Confidence Intervals

Brian Caffo

Department of Biostatistics
Johns Hopkins Bloomberg School of Public Health
Johns Hopkins University

August 23, 2012

Table of contents

- 1 Table of contents
- 2 Confidence intervals
- 3 CI for a normal variance
- 4 Student's t distribution
- 5 Confidence intervals for normal means
- 6 Profile likelihoods

Confidence intervals

- Previously, we discussed creating a confidence interval using the CLT
- Now we discuss the creation of better confidence intervals for small samples using Gosset's t distribution
- To discuss the t distribution we must discuss the Chi-squared distribution
- Throughout we use the following general procedure for creating CIs
 - a. Create a **Pivot** or statistic that does not depend on the parameter of interest
 - b. Solve the probability that the pivot lies between bounds for the parameter

The Chi-squared distribution

- Suppose that S^2 is the sample variance from a collection of iid $N(\mu, \sigma^2)$ data; then

$$\frac{(n-1)S^2}{\sigma^2} \sim \chi_{n-1}^2$$

which reads: follows a Chi-squared distribution with $n - 1$ degrees of freedom

- The Chi-squared distribution is skewed and has support on 0 to ∞
- The mean of the Chi-squared is its degrees of freedom
- The variance of the Chi-squared distribution is twice the degrees of freedom

Confidence interval for the variance

Note that if $\chi_{n-1,\alpha}^2$ is the α quantile of the Chi-squared distribution then

$$\begin{aligned} 1 - \alpha &= P\left(\chi_{n-1,\alpha/2}^2 \leq \frac{(n-1)S^2}{\sigma^2} \leq \chi_{n-1,1-\alpha/2}^2\right) \\ &= P\left(\frac{(n-1)S^2}{\chi_{n-1,1-\alpha/2}^2} \leq \sigma^2 \leq \frac{(n-1)S^2}{\chi_{n-1,\alpha/2}^2}\right) \end{aligned}$$

So that

$$\left[\frac{(n-1)S^2}{\chi_{n-1,1-\alpha/2}^2}, \frac{(n-1)S^2}{\chi_{n-1,\alpha/2}^2} \right]$$

is a $100(1 - \alpha)\%$ confidence interval for σ^2

Notes about this interval

- This interval relies heavily on the assumed normality
- Square-rooting the endpoints yields a CI for σ
- It turns out that

$$(n-1)S^2 \sim \text{Gamma}\{(n-1)/2, 2\sigma^2\}$$

which reads: follows a gamma distribution with shape $(n-1)/2$ and scale $2\sigma^2$

- Therefore, this can be used to plot a likelihood function for σ^2

Example


- A recent study 513 of organo-lead manufacturing workers reported an average total brain volume of $1,150.315 \text{ cm}^3$ with a standard deviation of 105.977. Assuming normality of the underlying measurements, calculate a confidence interval for the population variation in total brain volume.

Example continued

```
##CI for the variance
s2 <- 105.977 ^ 2
n <- 513
alpha <- .05
qtiles <- qchisq(c(alpha/2, 1 - alpha/2),
                  n - 1)
ival <- rev((n - 1) * s2 / qtiles)
##interval for the sd
sqrt(ival)
[1] 99.86484 112.89216
```


Plot the likelihood

```
sigmaVals <- seq(90, 120, length = 1000)
likeVals <- dgamma((n - 1) * s2,
                   shape = (n - 1)/2,
                   scale = 2*sigmaVals^2)
likeVals <- likeVals / max(likeVals)
plot(sigmaVals, likeVals, type = "l")
lines(range(sigmaVals[likeVals >= 1 / 8]),
      c(1 / 8, 1 / 8))
lines(range(sigmaVals[likeVals >= 1 / 16]),
      c(1 / 16, 1 / 16))
```



Gosset's t distribution

- Invented by William Gosset (under the pseudonym "Student") in 1908
- Has thicker tails than the normal
- Is indexed by a degrees of freedom; gets more like a standard normal as df gets larger
- Is obtained as

$$\frac{Z}{\sqrt{\frac{\chi^2}{df}}}$$

where Z and χ^2 are independent standard normals and Chi-squared distributions respectively

- Suppose that (X_1, \dots, X_n) are iid $N(\mu, \sigma^2)$, then:

a. $\frac{\bar{X} - \mu}{\sigma/\sqrt{n}}$ is standard normal

b. $\sqrt{\frac{(n-1)S^2}{\sigma^2(n-1)}} = S/\sigma$ is the square root of a Chi-squared divided by its df

- Therefore

$$\frac{\frac{\bar{X} - \mu}{\sigma/\sqrt{n}}}{S/\sigma} = \frac{\bar{X} - \mu}{S/\sqrt{n}}$$

follows Gosset's t distribution with $n - 1$ degrees of freedom

Confidence intervals for the mean

- Notice that the t statistic is a pivot, therefore we use it to create a confidence interval for μ
- Let $t_{df,\alpha}$ be the α^{th} quantile of the t distribution with df degrees of freedom

$$\begin{aligned} & 1 - \alpha \\ = & P \left(-t_{n-1, 1-\alpha/2} \leq \frac{\bar{X} - \mu}{S/\sqrt{n}} \leq t_{n-1, 1-\alpha/2} \right) \\ = & P \left(\bar{X} - t_{n-1, 1-\alpha/2} S/\sqrt{n} \leq \mu \leq \bar{X} + t_{n-1, 1-\alpha/2} S/\sqrt{n} \right) \end{aligned}$$

- Interval is $\bar{X} \pm t_{n-1, 1-\alpha/2} S/\sqrt{n}$

Note's about the t interval

- The t interval technically assumes that the data are iid normal, though it is robust to this assumption
- It works well whenever the distribution of the data is roughly symmetric and mound shaped
- Paired observations are often analyzed using the t interval by taking differences
- For large degrees of freedom, t quantiles become the same as standard normal quantiles; therefore this interval converges to the same interval as the CLT yielded

- For skewed distributions, the spirit of the t interval assumptions are violated
- Also, for skewed distributions, it doesn't make a lot of sense to center the interval at the mean
- In this case, consider taking logs or using a different summary like the median
- For highly discrete data, like binary, other intervals are available

Sleep data

Table of
contents

Confidence
intervals

CI for a
normal
variance

Student's t
distribution

Confidence
intervals for
normal means

Profile
likelihoods

In R typing `data(sleep)` brings up the sleep data originally analyzed in Gosset's Biometrika paper, which shows the increase in hours for 10 patients on two soporific drugs. R treats the data as two groups rather than paired.

Patient	g1	g2	diff
1	0.7	1.9	1.2
2	-1.6	0.8	2.4
3	-0.2	1.1	1.3
4	-1.2	0.1	1.3
5	-0.1	-0.1	0.0
6	3.4	4.4	1.0
7	3.7	5.5	1.8
8	0.8	1.6	0.8
9	0.0	4.6	4.6
10	2.0	3.4	1.4


```
data(sleep)
g1 <- sleep$extra[1 : 10]
g2 <- sleep$extra[11 : 20]
difference <- g2 - g1
mn <- mean(difference)#1.67
s <- sd(difference)#1.13
n <- 10
mn + c(-1, 1) * qt(.975, n-1) * s / sqrt(n)
t.test(difference)$conf.int
[1] 0.7001142 2.4598858
```

The non-central t distribution

- If X is $N(\mu, \sigma^2)$ and χ^2 is a Chi-squared random variable with df degrees of freedom then $\frac{X/\sigma}{\sqrt{\chi^2/df}}$ is called a **non-central t** random variable with non-centrality parameter μ/σ
- Note that
 - a. \bar{X} is $N(\mu, \sigma^2/n)$
 - b. $(n-1)S^2/\sigma^2$ is Chi-squared with $n-1$ df
- Then $\sqrt{n}\bar{X}/S$ is non-central t with non-centrality parameter $\sqrt{n}\mu/\sigma$
- We can use this to create a likelihood for μ/σ , the **effect size**

Starting after the code for the t interval

```
tStat <- sqrt(n) * mn / s
esVals <- seq(0, 1, length = 1000)
likVals <- dt(tStat, n - 1, ncp = sqrt(n) * esVals)
likVals <- likVals / max(likVals)
plot(esVals, likVals, type = "l")
lines(range(esVals[likVals>1/8]), c(1/8,1/8))
lines(range(esVals[likVals>1/16]), c(1/16,1/16))
```



The profile likelihood

- To obtain a likelihood for μ alone, the preferred method is called **profiling**
- The profile likelihood gets its name because the result is like the shadow you would get if you were to shine a light on the two-dimensional likelihood for μ and σ
- The profile likelihood for parameter value μ_0 is obtained by maximizing the joint likelihood for σ with μ fixed at μ_0
- This process is repeated for lots of values of μ_0

Calculating the profile likelihood

- The joint likelihood with μ fixed at μ_0 is

$$\begin{aligned} &\propto \prod_{i=1}^n \left[(\sigma^2)^{-1/2} \exp \left\{ -(x_i - \mu_0)^2 / 2\sigma^2 \right\} \right] \\ &= (\sigma^2)^{-n/2} \exp \left\{ - \sum_{i=1}^n (x_i - \mu_0)^2 / 2\sigma^2 \right\} \end{aligned}$$

- With μ_0 fixed, the maximum likelihood estimator for σ^2 is $\sum_{i=1}^n (x_i - \mu_0)^2 / n$ (homework)
- Plugging this back into the likelihood we get

$$\left(\sum_{i=1}^n (x_i - \mu_0)^2 / n \right)^{-n/2} \exp(-n/2)$$

- Therefore, removing multiplicative constants, the profile likelihood is

$$\left(\sum_{i=1}^n (x_i - \mu)^2 \right)^{-n/2}$$

- Note that this is clearly maximized at $\mu = \bar{X}$, the same as the ML estimate for μ for the complete likelihood

Some code

```
muVals <- seq(0, 3, length = 1000)
likVals <- sapply(muVals,
                  function(mu){
                    (sum((difference-mu)^2) /
                     sum((difference-mn)^2))^(n/2)
                  })
plot(muVals, likVals, type = "l")
lines(range(muVals[likVals>1/8]), c(1/8,1/8))
lines(range(muVals[likVals>1/16]), c(1/16,1/16))
```

