Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Mathematical Biostatistics Boot Camp: Lecture 11, Plotting

Brian Caffo

Department of Biostatistics
Johns Hopkins Bloomberg School of Public Health
Johns Hopkins University

September 19, 2012

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Table of
contents

Histograms

Stem and leaf

Dotcharts

Boxplots

KDEs

QQ-plots

Mosaic plots

Table of contents

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Histograms

- Histograms display a sample estimate of the density or mass function by plotting a bar graph of the frequency or proportion of times that a variable takes specific values, or a range of values for continuous data, within a sample

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Example

- The data set islands in the R package datasets contains the areas of all land masses in thousands of square miles
- Load the data set with the command data(islands)
- View the data by typing islands
- Create a histogram with the command hist(islands)
- Do ?hist for options

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Histogram of islands

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Pros and cons

- Histograms are useful and easy, apply to continuous, discrete and even unordered data
- They use a lot of ink and space to display very little information
- It's difficult to display several at the same time for comparisons

Also, for this data it's probably preferable to consider log base 10 (orders of magnitude), since the raw histogram simply says that most islands are small

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Histogram of log10(islands)

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Stem-and-leaf plots

- Stem-and-leaf plots are extremely useful for getting distribution information on the fly

- Read the text about creating them

- They display the complete data set and so waste very little ink

- Two data sets' stem and leaf plots can be shown back-to-back for comparisons

- Created by John Tukey, a leading figure in the development of the statistical sciences and signal processing

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Example

```
> stem(log10(islands))

  The decimal point is at the |

  1 | 1111112222233444
  1 | 5555556666667899999
  2 | 3344
  2 | 59
  3 |
  3 | 5678
  4 | 012
```

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Dotcharts

- Dotcharts simply display a data set, one point per dot
- Ordering of the of the dots and labeling of the axes can the display additional information
- Dotcharts show a complete data set and so have high data density
- May be impossible to construct/difficult to interpret for data sets with lots of points

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

islands data: log10(area) (log10(sq. miles))

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Discussion

- Maybe ordering alphabetically isn't the best thing for this data set
- Perhaps grouped by continent, then nations by geography (grouping Pacific islands together)?

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Dotplots comparing grouped data

- For data sets in groups, you often want to display density information by group

- If the size of the data permits, it displaying the whole data is preferable

- Add horizontal lines to depict means, medians

- Add vertical lines to depict variation, show confidence intervals interquartile ranges

- Jitter the points to avoid overplotting (jitter)

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Example

- The InsectSprays dataset contains counts of insect deaths by insecticide type (A, B, C, D, E, F)
- You can obtain the data set with the command

  `data(InsectSprays)`

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

The gist of the code is below

```
attach(InsectSprays)
plot(c(.5, 6.5), range(count))
sprayTypes <- unique(spray)
for (i in 1 : length(sprayTypes)){
  y <- count[spray == sprayTypes[i]]
  n <- sum(spray == sprayTypes[i])
  points(jitter(rep(i, n), amount = .1), y)
  lines(i + c(.12, .28), rep(mean(y), 2), lwd = 3)
  lines(rep(i + .2, 2),
        mean(y) + c(-1.96, 1.96) * sd(y) / sqrt(n)
        )
}
```

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Boxplots

- Boxplots are useful for the same sort of display as the dot chart, but in instances where displaying the whole data set is not possible

- Centerline of the boxes represents the median while the box edges correspond to the quartiles

- Whiskers extend out to a constant times the IQR or the max value

- Sometimes potential outliers are denoted by points beyond the whiskers

- Also invented by Tukey

- Skewness indicated by centerline being near one of the box edges

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Boxplots discussion

- Don't use boxplots for small numbers of observations, just plot the data!

- Try logging if some of the boxes are too squished relative to other ones; you can convert the axis to unlogged units (though they will not be equally spaced anymore)

- For data with lots and lots of observations omit the outliers plotting if you get so many of them that you cant see the points

- Example of a bad box plot

  `boxplot(rt(500, 2))`

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Kernel density estimates

- Kernel density estimates are essentially more modern versions of histograms providing density estimates for continuous data

- Observations are weighted according to a "kernel", in most cases a Gaussian density

- "Bandwidth" of the kernel effectively plays the role of the bin size for the histogram
  a. Too low of a bandwidth yields a too variable (jagged) measure of the density
  b. Too high of a bandwidth oversmooths

- The R function density can be used to create KDEs

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Example

Data is the waiting and eruption times in minutes between eruptions of the Old Faithful Geyser in Yellowstone National park

```
data(faithful)
d <- density(faithful$eruptions, bw = "sj")
plot(d)
```

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

N = 272   Bandwidth = 0.14

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Imaging example

- Consider the following image slice (created in R) from a high resolution MRI of a brain
- This is a single (axial) slice of a three-dimensional image
- Consider discarding the location information and plotting a KDE of the intensities

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

QQ-plots

- QQ-plots (for quantile-quantile) are extremely useful for comparing data to a theoretical distribution
- Plot the empirical quantiles against theoretical quantiles
- Most useful for diagnosing normality

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

- Let x_p be the p^{th} quantile from a $N(\mu, \sigma^2)$
- Then $P(X \leq x_p) = p$
- Clearly $P(Z \leq \frac{x_p - \mu}{\sigma}) = p$
- Therefore $x_p = \mu + z_p \sigma$ (this should not be news)
- Result, quantiles from a $N(\mu, \sigma^2)$ population should be linearly related to standard normal quantiles
- A normal qq-plot plot the empirical quantiles against the theoretical standard normal quantiles
- In R qqnorm for a normal QQ-plot and qqplot for a qqplot against an arbitrary distribution

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Normal Q−Q Plot

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Normal Q−Q Plot

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Normal Q–Q Plot

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Table of
contents

Histograms

Stem and leaf

Dotcharts

Boxplots

KDEs

QQ-plots

Mosaic plots

Mosaic plots

- Mosaic plots are useful for displaying contingency table data
- Consider Fisher's data regarding hair and eye color data for people from Caithness

```
library(MASS)
data(caith)
caith
mosaicplot(caith, color = topo.colors(4),
           main = "Mosiac plot")
```

	fair	red	medium	dark	black
blue	326	38	241	110	3
light	688	116	584	188	4
medium	343	84	909	412	26
dark	98	48	403	681	85

Mathematical
Biostatistics
Boot Camp:
Lecture 11,
Plotting

Brian Caffo

Mosiac plot